

COMPASS Architecture Docs — Consolidated Overview (one page)

A single at-a-glance index of every doc in `docs/architecture` + a current DB summary. Korean mirror: `INDEX_KR.md`. See each doc / `models.py` for detail. (Current schema: **15 app tables** + `alembic_version`, alembic head **029**.)

1. Docs in this folder (what each covers)

Doc	Lang	What it explains
<code>DATABASE_SCHEMA.md</code> / <code>_KR</code>	EN/KR	Table catalog — each of the 15 tables: purpose · why it exists · the anchor (parent-child) pattern
<code>DB_TABLES_ROLES.md</code> / <code>_KR</code>	EN/KR	One file's journey — transcript→NLP→AI→screen, told as a story of which drawer (table) gets what at each step
<code>AI_PIPELINE_NLP_DB_TABLES.md</code> / <code>_KR</code>	EN/KR	Pipeline detail — NLP 7 steps · AI 5 steps → table-fill matrix · ERD · per-table detail
<code>SAME_PATIENT_DIFF_DOCTOR.md</code> / <code>_KR</code>	EN/KR	Why same-patient/different-doctor data never mixes (file = patient+doctor+date is the key)
<code>ARCHITECTURE.md</code>	EN	System overview · deployment · module layout

2. Current DB — 15 app tables (by group)

Pipeline persistence (7) <code>transcript_analysis_log</code> (run header · anchor) · <code>sentence_prediction</code> (top-N sentences) · <code>nlp_all_predictions</code> (all sentences × 5 models) · <code>nlp_pipeline_intermediate</code> (NLP snapshots) · <code>llm_pipeline_intermediate</code> (AI intermediate) · <code>llm_domain_scoring_and_summary</code> (final, patient-facing) · <code>patient_summary</code> (patient parent anchor)
Behavior tracking (3) — UI interaction events (not values) <code>patient_report_page_behavior</code> (first-visit report page) · <code>patient_followup_survey_page_behavior</code> (survey pages: follow-up + first-visit Risk) · <code>doctor_behavior</code> (doctor dashboard)
Auth (2) <code>auth_user</code> · <code>auth_api_key</code>
Other (3) <code>session_recording</code> · <code>patient_survey_submission_log</code> (all patient survey answers — first-visit Risk = risk_perception_2 + follow-up sdm/dcs/satisfaction) · <code>doctor_rewrite_log</code> (doctor sentence rewrites)
System (1) <code>alembic_version</code> (migration version)

3. Three core concepts

- Anchor (parent-child):** `transcript_analysis_log` (id) and `patient_summary` (file, speaker) hold identity only; children FK to them → integrity · cascade · re-processing survival.
- Data flow:** file → NLP (sentences-probs ②③) → AI (scores-patient summaries ④⑤) → screens read ③⑤. Human input (surveys-rewrites) and behavior logs pile up separately.
- File = key:** records are distinguished by file (patient+doctor+date), not patient id → same-patient/different-doctor never mixes.

One-liner: all patient survey answers live in one place (`patient_survey_submission_log`), behavior logs in `*_page_behavior` / `doctor_behavior`, pipeline outputs under `transcript_analysis_log` 's children — 15 app tables total.